LOGSIGHT.ai

AI-driven log analytics

Discover the power of your logs

Over the next 5 years, the adoption of Microservices and Cloud computing will continuously increase

growth per year

18 %
Cloud Computing Market Size

21.7 %
Cloud Microservices Market

- IoT
- Industry 4.0
- Edge computing
- Self-driving cars
- Smart homes
- Smart cities

CRITICAL CLOUD INFRASTRUCTURES
Availability and **Reliability** are Key

solarwinds pingdom

State Of The Internet LIVE

Failures are inevitable

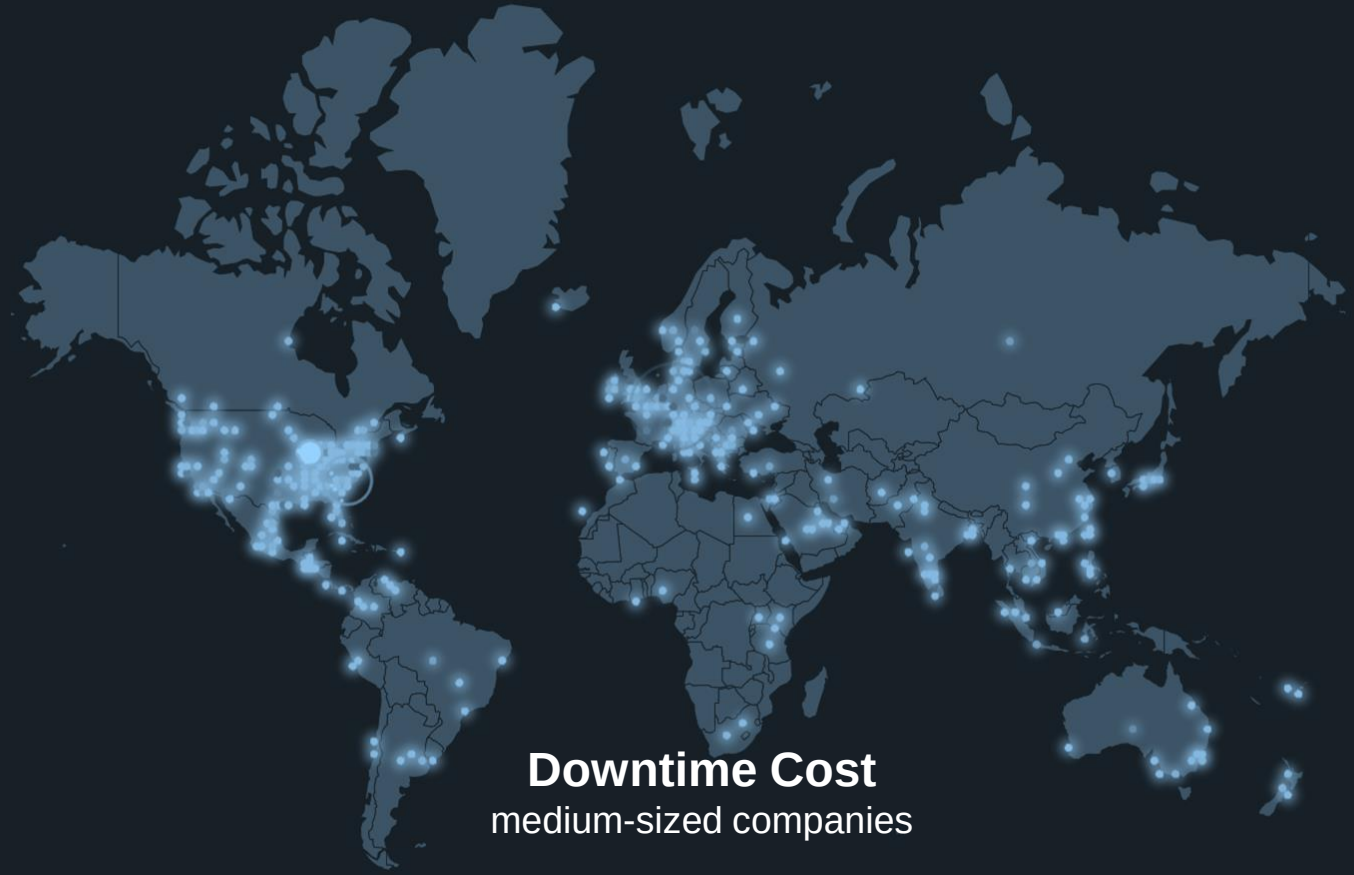

Downtime Cost

medium-sized companies

>$40.000/h

OUTAGES PER HOUR

0 +200

This is a live broadcast of Internet statistics, as collected by Pingdom from over 700,000 users across the world. See how you can use Pingdom to monitor your website.

Microsoft Office 365 Outage Reported For Some Users

The company says it has been 'applying steps to mitigate the issue' of connection timeouts for some mailbox users.

By Kyle Alspach — January 24, 2019, 12:37 PM EST

Slack
@SlackHQ

Replying to @sarbbottam

Sorry for the trouble! We'r
the moment, and hope to
normal as quickly as we ca

10:32 PM · Jun 28, 2019 · Sparkcentral.com

INTEGRATED TELECOMMUNICATIONS SERVICES MARCH 3, 2021 / 5:51 PM / UPDATED 19 DAYS AGO

Zoom down for more than a thousand users - Downdetector

By Reuters Staff — 1 MIN READ

Slack System Status
Resources for real-time Slack service.
status.slack.com

Microsoft Teams down: video call service back online

By Mike Moore 7 days ago

Microsoft Teams saw major outages, but should be back online now

Google Cloud Outage Attributed To 'Infrastructure Components' Issues

'We are experiencing an issue with Google Cloud infrastructure components,' Google Cloud confirmed in a post on its Google Cloud Status Dashboard.

By Donna Goodison — March 26, 2020, 12:55 PM EDT

Facebook struggles to deal with epic outage

By Donie O'Sullivan and Heather Kelly, CNN Business

Updated 0654 GMT (1454 HKT) March 14, 2019

amazon.com

Oops!

We're very sorry, but we're having trouble doing what you just asked us to do. Please give us another chance--click the Back button on your browser and try your request again. Or start from the beginning on our homepage.

TECH

Slack asks investors to trust that outage costs were a 'one-time' issue

PUBLISHED THU, SEP 5 2019•4:31 PM EDT

Jordan Novet
@JORDANNOVET

SHARE

Cloud

5-minute outage costs Google $545,000 in revenue

DYLAN TWENEY @DYLAN20 AUGUST 16, 2013 4:06 PM

Fact

DevOps use logs to repair IT failures

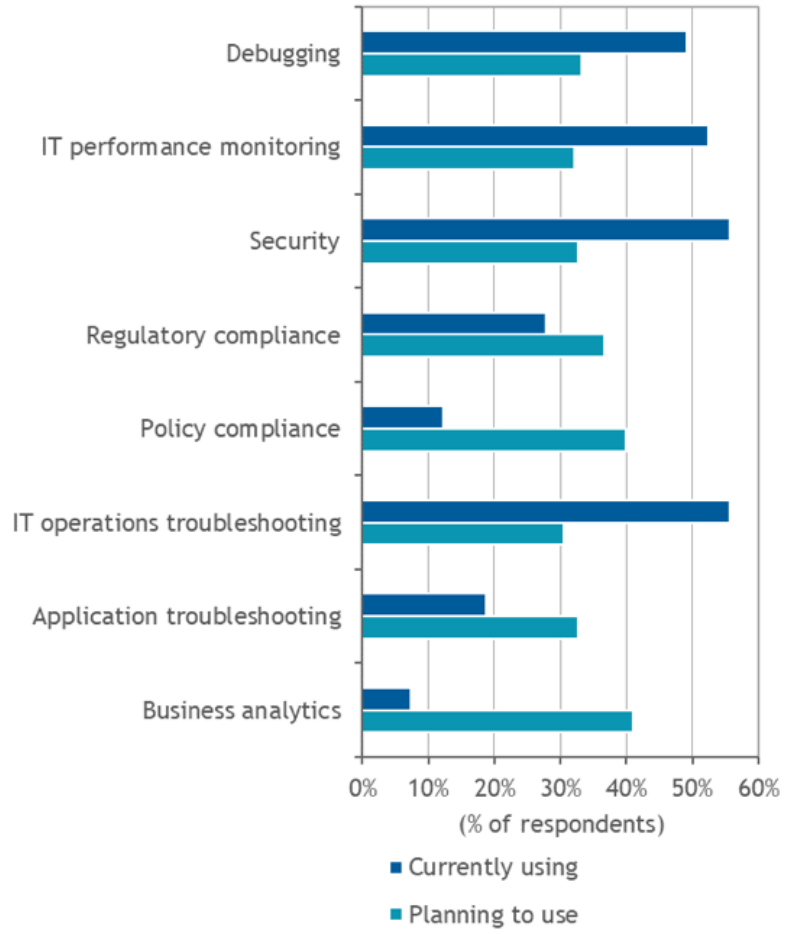

IDC, 2019

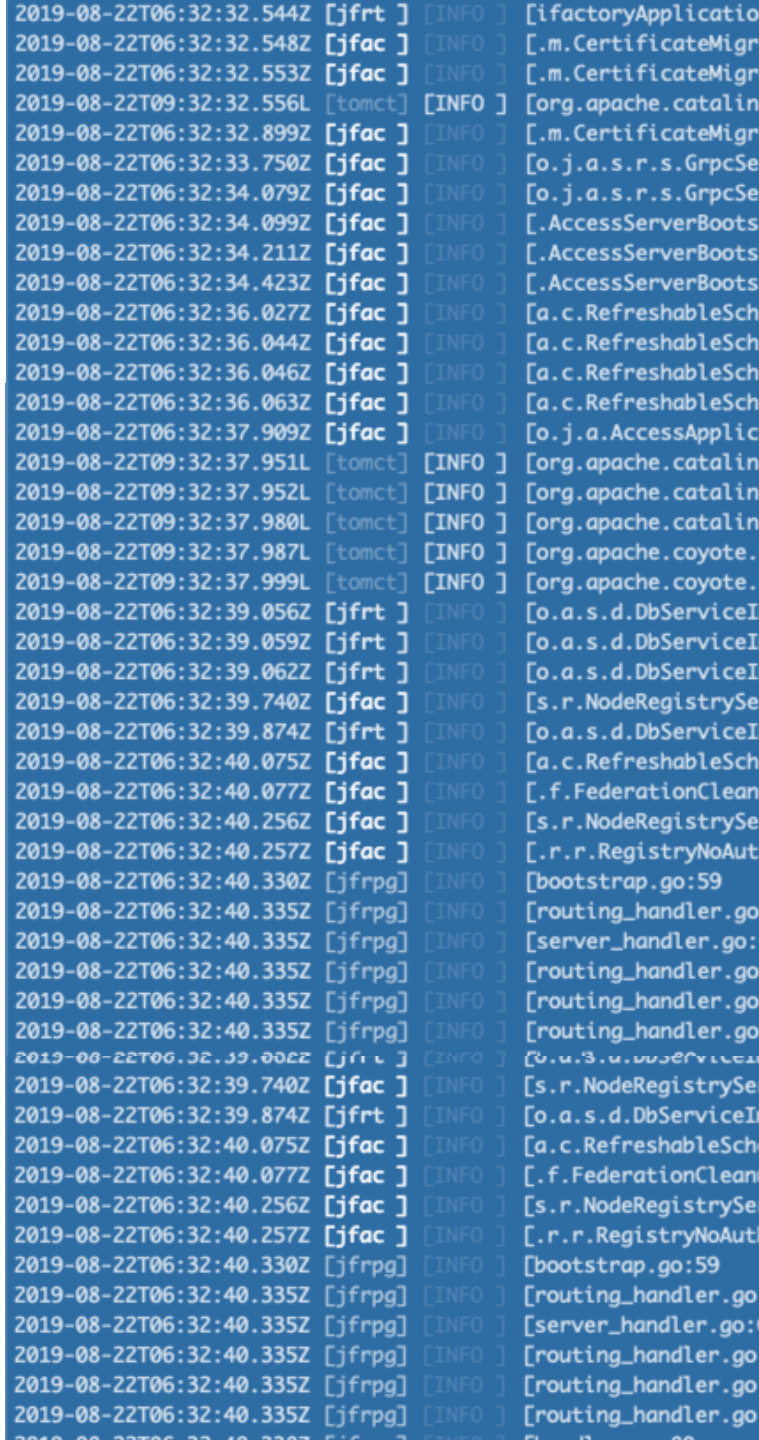

Problem

DevOps have two time-consuming options to troubleshoot IT failures

Log Search

(most companies choose this)

- Cheap, but requires a lot of manual work
- Does not scale to complex systems
- A lot of guesswork

Log Analysis

(fewer companies choose this)

- Simple, but expensive (requires many DevOps/SRE)
- High potential to overlook errors
- Human fatigue

Solution

Logsight.ai delivers automated log analytics through AI and pre-trained failure models

How It Works

We pre-train generic AI failure models

specialized with customers' systems behavior

We receive and process live logs from customers

notify customers if incidents happen

Customers receive detailed reports

Beyond the Status Quo
From log analysis to log analytics

How does it works

Log analysis
Clustering + Time Series

1. Unstructured

```
Dec 10 06:55:46: Invalid user webmaster from 173
Dec 10 06:55:46: input_userauth_request: invalid u
Dec 10 06:55:46: pam_unix(sshd:auth): check pas
Dec 10 06:55:48: Connection closed by 173.234.3
Dec 10 07:02:47: Connection closed by 212.47.254
Dec 10 07:07:38: Invalid user test9 from 52.80.34.
Dec 10 07:07:38: input_userauth_request: invalid u
Dec 10 07:07:38: pam_unix(sshd:auth): check pas
Dec 10 07:07:45: Failed password for invalid user
Dec 10 07:07:45: Received disconnect from 52.80.
Dec 10 07:08:28: Invalid user webmaster from 173
Dec 10 07:08:28: input_userauth_request: invalid u
Dec 10 07:08:28: pam_unix(sshd:auth): check pas
Dec 10 07:08:30: Connection closed by 173.234.3
Dec 10 07:11:42: Invalid user chen from 202.100.1
```

2. Clustering

```
Dec 10 06:55:46: Invalid user webmaster from 173
Dec 10 06:55:46: input_userauth_request: invalid u
Dec 10 06:55:46: pam_unix(sshd:auth): check pas
Dec 10 06:55:48: Connection closed by 173.234.3
Dec 10 07:02:47: Connection closed by 212.47.254
Dec 10 07:07:38: Invalid user test9 from 52.80.34.
Dec 10 07:07:38: input_userauth_request: invalid u
Dec 10 07:07:38: pam_unix(sshd:auth): check pas
Dec 10 07:07:45: Failed password for invalid user
Dec 10 07:07:49: Failed password for invalid user
Dec 10 07:08:15: Failed password for invalid user
Dec 10 07:08:28: input_userauth_request: invalid u
Dec 10 07:08:28: pam_unix(sshd:auth): check pas
Dec 10 07:08:30: Connection closed by 173.234.3
Dec 10 07:11:42: Invalid user chen from 202.100.1
```

3. Time series

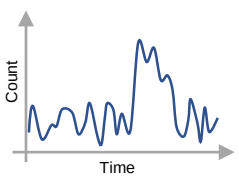

4. Anomaly detection

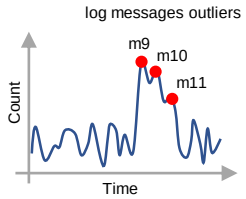

AI-driven log analytics
Deep Learning + Failure models

0. AI Model

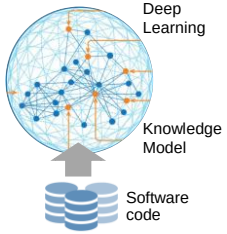

1. Unstructured

```
Dec 10 06:55:46: Invalid user webmaster from 173
Dec 10 06:55:46: input_userauth_request: invalid u
Dec 10 06:55:46: pam_unix(sshd:auth): check pas
Dec 10 06:55:48: Connection closed by 173.234.3
Dec 10 07:02:47: Connection closed by 212.47.254
Dec 10 07:07:38: Invalid user test9 from 52.80.34.
Dec 10 07:07:38: input_userauth_request: invalid u
Dec 10 07:07:38: pam_unix(sshd:auth): check pas
Dec 10 07:07:45: Failed password for invalid user
Dec 10 07:07:45: Received disconnect from 52.80.
Dec 10 07:08:28: Invalid user webmaster from 173
Dec 10 07:08:28: input_userauth_request: invalid u
Dec 10 07:08:28: pam_unix(sshd:auth): check pas
Dec 10 07:08:30: Connection closed by 173.234.3
Dec 10 07:11:42: Invalid user chen from 202.100.1
```

2. Structured

```
"httpRequest":{
  "requestMethod":"GET"
  },
"times":"2020-10-12T07:20:50.52Z",
"logging.googleapis.com/insertId":"42",
"logging.googleapis.com/labels":{
  "user_label_1":"value_1",
  "user_label_2":"value_2"
  },
"logging.googleapis.com/operation":{
  "id":"get_data",
  "producer":"github.com/MyProject/MyApplication",
  "first":"true"
  },
"logging.googleapis.com/sourceLocation":{
```

3. AI Inference

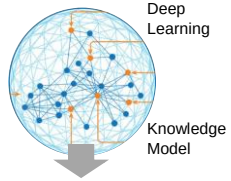

Invalid user webmaster from 173.234.31.186

Experience Report: System Log Analysis for Anomaly Detection, 2016

Use Case

Log search versus AI-driven log analytics

~350.000 log lines per day
60 types of errors

Log search

240 minutes

- Experienced DevOps conduct troubleshooting
- Reading logs and searching for hints
- Search filters: grep, awk … and error, warning …

AI-driven log analytics

4 minutes

- 59s uploading ~350.000 log lines
- 64s pre-processing
- 90s AI-driven log analytics

Beyond the Status Quo

From search, and analysis to log analytics

Log search
grep, awk, kql

Log analysis
Clustering + Time Series

Log analytics
Deep Learning + Failure models

Innovation

The three fundamental pillars

Leading edge AI technology

Learned from +100M code lines
Single and customized AI-models

Industry-ready

Deployed, tested, and continuously improved in collaboration with a large cloud provider

Long-term research

Results from years of research
5 PhD thesis
50+ publications


LOGSIGHT.ai
AI-driven log analytics
Discover the power of your logs